

Beyond Task Completion: A Rubric for Evaluating Agent-to-Agent Marketplaces

Ben Trovato*
G.K.M. Tobin*
trovato@corporation.com
webmaster@marysville-ohio.com
Institute for Clarity in Documentation
Dublin, Ohio, USA

Aparna Patel
Rajiv Gandhi University
Doimukh, Arunachal Pradesh, India

John Smith
The Thørväld Group
Hekla, Iceland
jsmith@affiliation.org

Lars Thørväld
The Thørväld Group
Hekla, Iceland
larst@affiliation.org

Huifen Chan
Tsinghua University
Haidian Qu, Beijing Shi, China

Julius P. Kumquat
The Kumquat Consortium
New York, USA
jpkumquat@consortium.net

Valerie Béranger
Inria Paris-Rocquencourt
Rocquencourt, France

Charles Palmer
Palmer Research Laboratories
San Antonio, Texas, USA
cpalmer@prl.com



Figure 1: Seattle Mariners at Spring Training, 2010.

Abstract

A clear and well-documented $\text{\LaTeX}$ document is presented as an article formatted for publication by ACM in a conference proceedings or journal publication. Based on the “acmart” document class, this article presents and explains many of the common variations, as well as many of the formatting elements an author may use in the preparation of the documentation of their work.

*Both authors contributed equally to this research.

Unpublished working draft. Not for distribution.

Permission to make digital or hard copies of all or part of this work for personal or professional use, or to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Conference acronym 'XX, Woodstock, NY
© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-1-4503-XXXX-X/2018/06
<https://doi.org/XXXXXXX.XXXXXXX>

CCS Concepts

• **Do Not Use This Code** → **Generate the Correct Terms for Your Paper**; *Generate the Correct Terms for Your Paper*; Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper.

Keywords

Do, Not, Use, This, Code, Put, the, Correct, Terms, for, Your, Paper

ACM Reference Format:

Ben Trovato, G.K.M. Tobin, Lars Thørväld, Valerie Béranger, Aparna Patel, Huifen Chan, Charles Palmer, John Smith, and Julius P. Kumquat. 2018. Beyond Task Completion: A Rubric for Evaluating Agent-to-Agent Marketplaces. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 9 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

ACM’s consolidated article template, introduced in 2017, provides a consistent L^AT_EX style for use across ACM publications, and incorporates accessibility and metadata-extraction functionality necessary for future Digital Library endeavors. Numerous ACM and SIG-specific L^AT_EX templates have been examined, and their unique features incorporated into this single new template.

If you are new to publishing with ACM, this document is a valuable guide to the process of preparing your work for publication. If you have published with ACM before, this document provides insight and instruction into more recent changes to the article template.

The “acmart” document class can be used to prepare articles for any ACM publication — conference or journal, and for any stage of publication, from review to final “camera-ready” copy, to the author’s own version, with very few changes to the source.

2 Template Overview

As noted in the introduction, the “acmart” document class can be used to prepare many different kinds of documentation — a double-anonymous initial submission of a full-length technical paper, a two-page SIGGRAPH Emerging Technologies abstract, a “camera-ready” journal article, a SIGCHI Extended Abstract, and more — all by selecting the appropriate *template style* and *template parameters*.

This document will explain the major features of the document class. For further information, the *L^AT_EX User’s Guide* is available from <https://www.acm.org/publications/proceedings-template>.

2.1 Template Styles

The primary parameter given to the “acmart” document class is the *template style* which corresponds to the kind of publication or SIG publishing the work. This parameter is enclosed in square brackets and is a part of the `documentclass` command:

```
\documentclass[STYLE]{acmart}
```

Journals use one of three template styles. All but three ACM journals use the `acmsmall` template style:

- `acmsmall`: The default journal template style.
- `acmlarge`: Used by JOCCH and TAP.
- `acmtog`: Used by TOG.

The majority of conference proceedings documentation will use the `acmconf` template style.

- `sigconf`: The default proceedings template style.
- `sigchi`: Used for SIGCHI conference articles.
- `sigplan`: Used for SIGPLAN conference articles.

2.2 Template Parameters

In addition to specifying the *template style* to be used in formatting your work, there are a number of *template parameters* which modify some part of the applied template style. A complete list of these parameters can be found in the *L^AT_EX User’s Guide*.

Frequently-used parameters, or combinations of parameters, include:

- `anonymous, review`: Suitable for a “double-anonymous” conference submission. Anonymizes the work and includes

line numbers. Use with the `\acmSubmissionID` command to print the submission’s unique ID on each page of the work.

- `authorversion`: Produces a version of the work suitable for posting by the author.
- `screen`: Produces colored hyperlinks.

This document uses the following string as the first command in the source file:

```
\documentclass[sigconf,authordraft]{acmart}
```

3 Modifications

Modifying the template — including but not limited to: adjusting margins, typeface sizes, line spacing, paragraph and list definitions, and the use of the `\vspace` command to manually adjust the vertical spacing between elements of your work — is not allowed.

Your document will be returned to you for revision if modifications are discovered.

4 Typefaces

The “acmart” document class requires the use of the “Libertine” typeface family. Your T_EX installation should include this set of packages. Please do not substitute other typefaces. The “lmodern” and “l^AT_EX” packages should not be used, as they will override the built-in typeface families.

5 Title Information

The title of your work should use capital letters appropriately — <https://capitalizemytitle.com/> has useful rules for capitalization. Use the `title` command to define the title of your work. If your work has a subtitle, define it with the `subtitle` command. Do not insert line breaks in your title.

If your title is lengthy, you must define a short version to be used in the page headers, to prevent overlapping text. The `title` command has a “short title” parameter:

```
\title[short title]{full title}
```

6 Authors and Affiliations

Each author must be defined separately for accurate metadata identification. As an exception, multiple authors may share one affiliation. Authors’ names should not be abbreviated; use full first names wherever possible. Include authors’ e-mail addresses whenever possible.

Grouping authors’ names or e-mail addresses, or providing an “e-mail alias,” as shown below, is not acceptable:

```
\author{Brooke Aster, David Mehldau}
\email{dave,judy,steve@university.edu}
\email{firstname.lastname@phillips.org}
```

The `authornote` and `authornotemark` commands allow a note to apply to multiple authors — for example, if the first two authors of an article contributed equally to the work.

If your author list is lengthy, you must define a shortened version of the list of authors to be used in the page headers, to prevent overlapping text. The following command should be placed just after the last `\author{}` definition:

```
\renewcommand{\shortauthors}{McCartney, et al.}
```

Omitting this command will force the use of a concatenated list of all of the authors' names, which may result in overlapping text in the page headers.

The article template's documentation, available at <https://www.acm.org/publications/proceedings-template>, has a complete explanation of these commands and tips for their effective use.

Note that authors' addresses are mandatory for journal articles.

7 Rights Information

Authors of any work published by ACM will need to complete a rights form. Depending on the kind of work, and the rights management choice made by the author, this may be copyright transfer, permission, license, or an OA (open access) agreement.

Regardless of the rights management choice, the author will receive a copy of the completed rights form once it has been submitted. This form contains L^AT_EX commands that must be copied into the source document. When the document source is compiled, these commands and their parameters add formatted text to several areas of the final document:

- the “ACM Reference Format” text on the first page.
- the “rights management” text on the first page.
- the conference information in the page header(s).

Rights information is unique to the work; if you are preparing several works for an event, make sure to use the correct set of commands with each of the works.

The ACM Reference Format text is required for all articles over one page in length, and is optional for one-page articles (abstracts).

8 CCS Concepts and User-Defined Keywords

Two elements of the “acmart” document class provide powerful taxonomic tools for you to help readers find your work in an online search.

The ACM Computing Classification System — <https://www.acm.org/publications/class-2012> — is a set of classifiers and concepts that describe the computing discipline. Authors can select entries from this classification system, via <https://dl.acm.org/ccs/ccs.cfm>, and generate the commands to be included in the L^AT_EX source.

User-defined keywords are a comma-separated list of words and phrases of the authors' choosing, providing a more flexible way of describing the research being presented.

CCS concepts and user-defined keywords are required for for all articles over two pages in length, and are optional for one- and two-page articles (or abstracts).

9 Sectioning Commands

Your work should use standard L^AT_EX sectioning commands: `\section`, `\subsection`, `\subsubsection`, `\paragraph`, and `\subparagraph`. The sectioning levels up to `\subsubsection` should be numbered; do not remove the numbering from the commands.

Simulating a sectioning command by setting the first word or words of a paragraph in boldface or italicized text is **not allowed**.

Below are examples of sectioning commands.

9.1 Subsection

This is a subsection.

Table 1: Frequency of Special Characters

Non-English or Math	Frequency	Comments
∅	1 in 1,000	For Swedish names
π	1 in 5	Common in math
\$	4 in 5	Used in business
Ψ ₁ ²	1 in 40,000	Unexplained usage

9.1.1 Subsubsection. This is a subsubsection.

Paragraph. This is a paragraph.
Subparagraph This is a subparagraph.

10 Tables

The “acmart” document class includes the “booktabs” package — <https://ctan.org/pkg/booktabs> — for preparing high-quality tables. Table captions are placed *above* the table.

Because tables cannot be split across pages, the best placement for them is typically the top of the page nearest their initial cite. To ensure this proper “floating” placement of tables, use the environment **table** to enclose the table's contents and the table caption. The contents of the table itself must go in the **tabular** environment, to be aligned properly in rows and columns, with the desired horizontal and vertical rules. Again, detailed instructions on **tabular** material are found in the *L^AT_EX User's Guide*.

Immediately following this sentence is the point at which Table 1 is included in the input file; compare the placement of the table here with the table in the printed output of this document.

To set a wider table, which takes up the whole width of the page's live area, use the environment **table*** to enclose the table's contents and the table caption. As with a single-column table, this wide table will “float” to a location deemed more desirable. Immediately following this sentence is the point at which Table 2 is included in the input file; again, it is instructive to compare the placement of the table here with the table in the printed output of this document.

Always use `midrule` to separate table header rows from data rows, and use it only for this purpose. This enables assistive technologies to recognise table headers and support their users in navigating tables more easily.

11 Math Equations

You may want to display math equations in three distinct styles: inline, numbered or non-numbered display. Each of the three are discussed in the next sections.

11.1 Inline (In-text) Equations

A formula that appears in the running text is called an inline or in-text formula. It is produced by the **math** environment, which can be invoked with the usual `\begin . . . \end` construction or with the short form `$ . . . $`. You can use any of the symbols and structures, from α to ω , available in L^AT_EX [29]; this section will simply show a few examples of in-text equations in context. Notice how this equation: $\lim_{n \rightarrow \infty} x = 0$, set here in in-line math style, looks slightly different when set in display style. (See next section).

Table 2: Some Typical Commands

Command	A Number	Comments
<code>\author</code>	100	Author
<code>\table</code>	300	For tables
<code>\table*</code>	400	For wider tables

11.2 Display Equations

A numbered display equation—one set off by vertical space from the text and centered horizontally—is produced by the `equation` environment. An unnumbered display equation is produced by the `displaymath` environment.

Again, in either environment, you can use any of the symbols and structures available in $\text{\LaTeX}$; this section will just give a couple of examples of display equations in context. First, consider the equation, shown as an inline equation above:

$$\lim_{n \rightarrow \infty} x = 0 \tag{1}$$

Notice how it is formatted somewhat differently in the `displaymath` environment. Now, we'll enter an unnumbered equation:

$$\sum_{i=0}^{\infty} x + 1$$

and follow it with another numbered equation:

$$\sum_{i=0}^{\infty} x_i = \int_0^{\pi+2} f \tag{2}$$

just to demonstrate $\text{\LaTeX}$'s able handling of numbering.

12 Figures

The “figure” environment should be used for figures. One or more images can be placed within a figure. If your figure contains third-party material, you must clearly identify it as such, as shown in the example below.

Your figures should contain a caption which describes the figure to the reader.

Figure captions are placed *below* the figure.

Every figure should also have a figure description unless it is purely decorative. These descriptions convey what's in the image to someone who cannot see it. They are also used by search engine crawlers for indexing images, and when images cannot be loaded.

A figure description must be unformatted plain text less than 2000 characters long (including spaces). **Figure descriptions should not repeat the figure caption – their purpose is to capture important information that is not already provided in the caption or the main text of the paper.** For figures that convey important and complex new information, a short text description may not be adequate. More complex alternative descriptions can be placed in an appendix and referenced in a short figure description. For example, provide a data table capturing the information in a bar chart, or a structured list representing a graph. For additional information regarding how best to write figure descriptions and why doing this is so important, please see <https://www.acm.org/publications/taps/describing-figures/>.



Figure 2: 1907 Franklin Model D roadster. Photograph by Harris & Ewing, Inc. [Public domain], via Wikimedia Commons. (<https://goo.gl/VLCRBB>).

12.1 The “Teaser Figure”

A “teaser figure” is an image, or set of images in one figure, that are placed after all author and affiliation information, and before the body of the article, spanning the page. If you wish to have such a figure in your article, place the command immediately before the `\maketitle` command:

```
\begin{teaserfigure}
  \includegraphics[width=\textwidth]{sampleteaser}
  \caption{figure caption}
  \Description{figure description}
\end{teaserfigure}
```

13 Citations and Bibliographies

The use of $\text{\BibTeX}$ for the preparation and formatting of one's references is strongly recommended. Authors' names should be complete — use full first names (“Donald E. Knuth”) not initials (“D. E. Knuth”) — and the salient identifying features of a reference should be included: title, year, volume, number, pages, article DOI, etc.

The bibliography is included in your source document with these two commands, placed just before the `\end{document}` command:

```
\bibliographystyle{ACM-Reference-Format}
\bibliography{bibfile}
```

where “bibfile” is the name, without the “.bib” suffix, of the $\text{\BibTeX}$ file.

Citations and references are numbered by default. A small number of ACM publications have citations and references formatted in the “author year” style; for these exceptions, please include this command in the **preamble** (before the command “\begin{document}”) of your $\text{\LaTeX}$ source:

```
\citestyle{acmauthoryear}
```

Some examples. A paginated journal article [2], an enumerated journal article [15], a reference to an entire issue [14], a monograph (whole book) [28], a monograph/whole book in a series (see 2a in spec. document) [22], a divisible-book such as an anthology or compilation [17] followed by the same example, however we only output the series if the volume number is given [18] (so Editor00a’s series should NOT be present since it has no vol. no.), a chapter in a divisible book [43], a chapter in a divisible book in a series [16], a multi-volume work as book [27], a couple of articles in a proceedings (of a conference, symposium, workshop for example) (paginated proceedings article) [4, 20], a proceedings article with all possible elements [42], an example of an enumerated proceedings article [19], an informally published work [21], a couple of preprints [9, 12], a doctoral dissertation [13], a master’s thesis: [5], an online document / world wide web resource [1, 35, 44], a video game (Case 1) [34] and (Case 2) [32] and [30] and (Case 3) a patent [41], work accepted for publication [38], ‘YYYYb’-test for prolific author [39] and [40]. Other cites might contain ‘duplicate’ DOI and URLs (some SIAM articles) [26]. Boris / Barbara Beeton: multi-volume works as books [24] and [23]. A presentation [37]. An article under review [11]. A couple of citations with DOIs: [25, 26]. Online citations: [44–46]. Artifacts: [36] and [8].

14 Acknowledgments

Identification of funding sources and other support, and thanks to individuals and groups that assisted in the research and the preparation of the work should be included in an acknowledgment section, which is placed just before the reference section in your document.

This section has a special environment:

```
\begin{acks}
...
\end{acks}
```

so that the information contained therein can be more easily collected during the article metadata extraction phase, and to ensure consistency in the spelling of the section heading.

Authors should not prepare this section as a numbered or un-numbered \section; please use the “acks” environment.

15 Appendices

If your work needs an appendix, add it before the “\end{document}” command at the conclusion of your source document.

Start the appendix with the “appendix” command:

```
\appendix
```

and note that in the appendix, sections are lettered, not numbered. This document has two appendices, demonstrating the section and subsection identification method.

16 Multi-language papers

Papers may be written in languages other than English or include titles, subtitles, keywords and abstracts in different languages (as a rule, a paper in a language other than English should include an English title and an English abstract). Use `language=...` for every language used in the paper. The last language indicated is the main language of the paper. For example, a French paper with additional titles and abstracts in English and German may start with the following command

```
\documentclass[sigconf, language=english, language=german,
language=french]{acmart}
```

The title, subtitle, keywords and abstract will be typeset in the main language of the paper. The commands `\translatedXXX`, `XXX` begin title, subtitle and keywords, can be used to set these elements in the other languages. The environment `translatedabstract` is used to set the translation of the abstract. These commands and environment have a mandatory first argument: the language of the second argument. See `sample-sigconf-i13n.tex` file for examples of their usage.

17 SIGCHI Extended Abstracts

The “sigchi-a” template style (available only in $\text{\LaTeX}$ and not in Word) produces a landscape-orientation formatted article, with a wide left margin. Three environments are available for use with the “sigchi-a” template style, and produce formatted output in the margin:

sidebar: Place formatted text in the margin.
marginfigure: Place a figure in the margin.
marginable: Place a table in the margin.

Acknowledgments

To Robert, for the bagels and explaining CMYK and color spaces.

References

- [1] Rafal Ablamowicz and Bertfried Fauser. 2007. *CLIFFORD: a Maple 11 Package for Clifford Algebra Computations, version 11*. Retrieved February 28, 2008 from <http://math.tntech.edu/rafal/cliff11/index.html>
- [2] Patricia S. Abril and Robert Plant. 2007. The patent holder’s dilemma: Buy, sell, or troll? *Commun. ACM* 50, 1 (Jan. 2007), 36–44. doi:10.1145/1188913.1188915
- [3] Abdulaziz Alzuabi, Yara Rizk, Tamar Lavee, Bhavya Chopra, Sathish Kumar Selvam, Keerthiram Murugesan, and Subhajit Chaudhury. 2024. NegotiationArena: Evaluating Neural Networks in Multi-Turn Bilateral Negotiation. arXiv:2402.05863 [cs.AI]
- [4] Sten Andler. 1979. Predicate Path expressions. In *Proceedings of the 6th. ACM SIGACT-SIGPLAN symposium on Principles of Programming Languages (POPL ’79)*. ACM Press, New York, NY, 226–236. doi:10.1145/567752.567774
- [5] David A. Anisi. 2003. *Optimal Motion Control of a Ground Vehicle*. Master’s thesis. Royal Institute of Technology (KTH), Stockholm, Sweden.
- [6] Anonymous. 2026. MarketBench: Agent Calibration and Auction Simulation. arXiv:2604.23897 [cs.AI]
- [7] Anthropic. 2025. Project Deal: Autonomous Agent Commerce with 69 Employees. Anthropic Internal Report. 69 Claude agents negotiated real goods in a shared Slack channel.
- [8] Sam Anzaroot and Andrew McCallum. 2013. *UMass Citation Field Extraction Dataset*. Retrieved May 27, 2019 from <http://www.iesl.cs.umass.edu/data/data-umasscitationfield>
- [9] Sam Anzaroot, Alexandre Passos, David Belanger, and Andrew McCallum. 2014. *Learning Soft Linear Constraints with Application to Citation Field Extraction*. arXiv:1403.1349 doi:10.48550/arXiv.1403.1349
- [10] Amos Azaria, Tomer Rokach, and Ariel Felner. 2026. When Agent Markets Arrive: Diagon and the GDP Multiplier. arXiv:2604.06688 [cs.MA]
- [11] R. Baggett, M. Simecek, C. Chambellan, K. Tsui, and M. Fraune. 2025. Fluidity in the Phased Framework of Technology Acceptance: Case Study to Gain a Holistic

- Understanding of (Older Adult) Participant Advancement Through Acceptance Phases with Mobile Telepresence Robots. *Robotics Aut. Systems*. Manuscript submitted for review.
- [12] Lutz Bornmann, K. Brad Wray, and Robin Haunschild. 2019. *Citation concept analysis (CCA)—A new form of citation analysis revealing the usefulness of concepts for other researchers illustrated by two exemplary case studies including classic books by Thomas S. Kuhn and Karl R. Popper*. arXiv:1905.12410 [cs.DL]
- [13] Kenneth L. Clarkson. 1985. *Algorithms for Closest-Point Problems (Computational Geometry)*. Ph. D. Dissertation. Stanford University, Palo Alto, CA. UMI Order Number: AAT 8506171.
- [14] Jacques Cohen (Ed.). 1996. Special issue: Digital Libraries. *Commun. ACM* 39, 11 (Nov. 1996).
- [15] Sarah Cohen, Werner Nutt, and Yehoshua Sagie. 2007. Deciding equivalences among conjunctive aggregate queries. *J. ACM* 54, 2, Article 5 (April 2007), 50 pages. doi:10.1145/1219092.1219093
- [16] Bruce P. Douglass, David Harel, and Mark B. Trakhtenbrot. 1998. Statecharts in use: structured analysis and object-orientation. In *Lectures on Embedded Systems*, Grzegorz Rozenberg and Frits W. Vaandrager (Eds.). Lecture Notes in Computer Science, Vol. 1494. Springer-Verlag, London, 368–394. doi:10.1007/3-540-65193-4_29
- [17] Ian Editor (Ed.). 2007. *The title of book one* (1st. ed.). The name of the series one, Vol. 9. University of Chicago Press, Chicago, Chapter The title of the chapter, 127–238. doi:10.1007/3-540-09237-4
- [18] Ian Editor (Ed.). 2008. *The title of book two* (2nd. ed.). University of Chicago Press, Chicago, Chapter 100, 25–137. doi:10.1007/3-540-09237-4
- [19] Matthew Van Gundy, Davide Balzarotti, and Giovanni Vigna. 2007. Catch me, if you can: Evading network signatures with web-based polymorphic worms. In *Proceedings of the first USENIX workshop on Offensive Technologies (WOOT '07)*. USENIX Association, Berkeley, CA, Article 7, 9 pages.
- [20] Torben Hagerup, Kurt Mehlhorn, and J. Ian Munro. 1993. Maintaining Discrete Probability Distributions Optimally. In *Proceedings of the 20th International Colloquium on Automata, Languages and Programming (Lecture Notes in Computer Science, Vol. 700)*. Springer-Verlag, Berlin, 253–264.
- [21] David Harel. 1978. *LOGICS of Programs: AXIOMATICS and DESCRIPTIVE POWER*. MIT Research Lab Technical Report TR-200. Massachusetts Institute of Technology, Cambridge, MA.
- [22] David Harel. 1979. *First-Order Dynamic Logic*. Lecture Notes in Computer Science, Vol. 68. Springer-Verlag, New York, NY. doi:10.1007/3-540-09237-4
- [23] Lars Hörmander. 1985. *The analysis of linear partial differential operators. III*. Grundlehren der Mathematischen Wissenschaften [Fundamental Principles of Mathematical Sciences], Vol. 275. Springer-Verlag, Berlin, Germany. viii+525 pages. Pseudodifferential operators.
- [24] Lars Hörmander. 1985. *The analysis of linear partial differential operators. IV*. Grundlehren der Mathematischen Wissenschaften [Fundamental Principles of Mathematical Sciences], Vol. 275. Springer-Verlag, Berlin, Germany. vii+352 pages. Fourier integral operators.
- [25] IEEE. 2004. IEEE TCSC Executive Committee. In *Proceedings of the IEEE International Conference on Web Services (ICWS '04)*. IEEE Computer Society, Washington, DC, USA, 21–22. doi:10.1109/ICWS.2004.64
- [26] Markus Kirschmer and John Voight. 2010. Algorithmic Enumeration of Ideal Classes for Quaternion Orders. *SIAM J. Comput.* 39, 5 (Jan. 2010), 1714–1747. doi:10.1137/080734467
- [27] Donald E. Knuth. 1997. *The Art of Computer Programming, Vol. 1: Fundamental Algorithms (3rd. ed.)*. Addison Wesley Longman Publishing Co., Inc., Boston.
- [28] David Kosiur. 2001. *Understanding Policy-Based Networking* (2nd. ed.). Wiley, New York, NY.
- [29] Leslie Lamport. 1986. *LaTeX: A Document Preparation System*. Addison-Wesley, Reading, MA.
- [30] Newton Lee. 2005. Interview with Bill Kinder: January 13, 2005. Video. *Comput. Entertain.* 3, 1, Article 4 (Jan.-March 2005). doi:10.1145/1057270.1057278
- [31] Ziwei Liu, Ping Luo, Shi Qiu, Xiaogang Wang, and Xiaoou Tang. 2016. DeepFashion: Powering Robust Clothes Recognition and Retrieval with Rich Annotations. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 1096–1104.
- [32] Dave Novak. 2003. Solder man. Video. In *ACM SIGGRAPH 2003 Video Review on Animation theater Program: Part I - Vol. 145 (July 27–27, 2003)*. ACM Press, New York, NY, 4. doi:10.945/woot07-S422 <http://video.google.com/videoplay?docid=6528042696351994555>
- [33] NVIDIA. 2025. NeMo Gym: Open-Source Agentic Evaluation Framework. <https://docs.nvidia.com/nemo/gym>. Accessed May 2026.
- [34] Barack Obama. 2008. A more perfect union. Video. Retrieved March 21, 2008 from <http://video.google.com/videoplay?docid=6528042696351994555>
- [35] Poker-Edge.Com. 2006. Stats and Analysis. Retrieved June 7, 2006 from <http://www.poker-edge.com/stats.php>
- [36] R Core Team. 2019. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. <https://www.R-project.org/>
- [37] Brian J. Reiser. 2014. Designing coherent storylines aligned with NGSS for the K-12 classroom. Presentation at National Science Education Leadership Association Meeting, Boston, MA, USA. <https://www.academia.edu/6884962/>
- [38] Bernard Rous. 2008. The Enabling of Digital Libraries. *Digital Libraries* 12, 3, Article 5 (July 2008). To appear.
- [39] Mehdi Saeedi, Morteza Saheb Zamani, and Mehdi Sedighi. 2010. A library-based synthesis methodology for reversible logic. *Microelectron. J.* 41, 4 (April 2010), 185–194.
- [40] Mehdi Saeedi, Morteza Saheb Zamani, Mehdi Sedighi, and Zahra Sasanian. 2010. Synthesis of Reversible Circuit Using Cycle-Based Approach. *J. Emerg. Technol. Comput. Syst.* 6, 4 (Dec. 2010), 12 pages.
- [41] Joseph Scientist. 2009. The fountain of youth. Patent No. 12345, Filed July 1st., 2008, Issued Aug. 9th., 2009.
- [42] Stan W. Smith. 2010. An experiment in bibliographic mark-up: Parsing metadata for XML export. In *Proceedings of the 3rd. annual workshop on Librarians and Computers (LAC '10, Vol. 3)*, Reginald N. Smythe and Alexander Noble (Eds.). Paparazzi Press, Milan Italy, 422–431.
- [43] Asad Z. Spector. 1990. Achieving application requirements. In *Distributed Systems* (2nd. ed.), Sape Mullender (Ed.). ACM Press, New York, NY, 19–33. doi:10.1145/90417.90738
- [44] Harry Thornburg. 2001. *Introduction to Bayesian Statistics*. Retrieved March 2, 2005 from <http://ccrma.stanford.edu/~jos/bayes/bayes.html>, archived at <https://web.archive.org/web/20240505055615/https://ccrma.stanford.edu/~jos/bayes/bayes.html>
- [45] TUG. 2017. *Institutional members of the T_EX Users Group*. Retrieved May 27, 2017 from <http://www.tug.org/instmem.html>
- [46] Boris Veytsman. 2017. *acmart—Class for typesetting publications of ACM*. Retrieved May 27, 2017 from <http://www.ctan.org/pkg/acmart>

A Research Methods

A.1 Part One

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Morbi malesuada, quam in pulvinar varius, metus nunc fermentum urna, id sollicitudin purus odio sit amet enim. Aliquam ullamcorper eu ipsum vel mollis. Curabitur quis dictum nisl. Phasellus vel semper risus, et lacinia dolor. Integer ultricies commodo sem nec semper.

A.2 Part Two

Etiam commodo feugiat nisl pulvinar pellentesque. Etiam auctor sodales ligula, non varius nibh pulvinar semper. Suspendisse nec lectus non ipsum convallis congue hendrerit vitae sapien. Donec at laoreet eros. Vivamus non purus placerat, scelerisque diam eu, cursus ante. Etiam aliquam tortor auctor efficitur mattis.

B Online Resources

Nam id fermentum dui. Suspendisse sagittis tortor a nulla mollis, in pulvinar ex pretium. Sed interdum orci quis metus euismod, et sagittis enim maximus. Vestibulum gravida massa ut felis suscipit congue. Quisque mattis elit a risus ultrices commodo venenatis eget dui. Etiam sagittis eleifend elementum.

Nam interdum magna at lectus dignissim, ac dignissim lorem rhoncus. Maecenas eu arcu ac neque placerat aliquam. Nunc pulvinar massa et mattis lacinia.

C Methodology

To study how market mechanics shape agent negotiation behavior, we design a controlled multi-agent simulation in which the same ten synthetic agents interact across two progressively complex market scenarios. The same agent personas, model assignments, and random seeds are reused across all three experimental conditions, so that any difference in observed behavior can be attributed to changes in the market rules alone — not to variation in agent identity or run initialisation.

Each condition is evaluated under four model configurations: a Sonnet focal agent against a Haiku field (*S* vs. *H*), a Haiku focal agent against a Sonnet field (*H* vs. *S*), all-Sonnet (*S* vs. *S*), and all-Haiku (*H* vs. *H*), using Claude Sonnet 4.5 and Haiku 4.5 via OpenRouter. Each condition runs across five persona sets and four model configurations, yielding 20 rollouts per condition and 60 rollouts in total.

C.1 Persona Sets

Each set contains ten agents, every one of whom simultaneously occupies the role of buyer and seller: a list of items to sell with private floor prices, a list of wants with private ceiling prices, a declared negotiating style, and a personal context block the agent is instructed not to disclose. Listing 1 shows a representative example. Private reservation values are embedded only in each agent’s own system prompt, establishing an information asymmetry that mirrors real negotiation conditions.

```
{
  "name": "Zara",
  "items_to_sell": [
    { "name": "Sony WH-CH520 headphones",
      "floor_price": 30 },
    { "name": "The Flavor Bible (hardcover)",
      "floor_price": 18 }
  ],
  "items_to_buy": [
    { "description": "Yoga mat or exercise mat",
      "ceiling_price": 25 }
  ],
  "style": "Playful, haggles cheerfully.",
  "private": {
    "occupation": "Freelance graphic designer",
    "financial_situation": "Car repair bill ($1,200).",
    "debt_context": "$3,400 in credit card debt."
  }
}
```

Listing 1: Representative agent persona (Zara, Set 03). The private block is injected only into the agent’s own system prompt.

The five sets span a deliberate range of structural market conditions — from configurations where every matching pair has overlapping price ranges and all deals are readily achievable, to sets containing price gaps that make certain trades impossible regardless of negotiation skill, multi-buyer competition for a single scarce item, and chain-structured dependencies in which one deal must close before another becomes viable. This diversity ensures that results reflect negotiation behavior across varied market structures rather than any single configuration.

C.2 Simulation Infrastructure

The simulation runs inside NVIDIA NeMo Gym [33], an open-source agentic evaluation framework that manages task dispatch, server lifecycle, and rollout collection at scale. The marketplace is implemented as a NeMo Gym *Resources Server* — a FastAPI application exposing the scenario’s tool endpoints and a */verify* endpoint — into which the focal agent issues sequential tool calls; each call is validated, written to the shared append-only channel, followed by one opponent response turn, and the updated state is

returned to the focal agent. Runs proceed for up to approximately 100 turns, after which */verify* scores all six rubric dimensions from the complete channel transcript and closed-deal ledger.

C.3 Experimental Scenarios

Scenario 1 — MarketDeal. Agents negotiate over listed items using monetary offers. Scenario 1 is evaluated under two market conditions. In the *baseline condition* (Stage I), each agent has access to six tools: *post_listing*, *make_offer*, *counter_offer*, *accept_offer*, *reject_offer*, and *pass*. When an offer is accepted, the ledger records the agreed price alongside each party’s private reservation value, enabling post-hoc surplus analysis. In the *reputation-augmented condition* (Stage II), a seventh tool — *lookup_agent* — allows the focal agent to retrieve a peer’s star rating and free-text reviews silently, without advancing the turn counter or broadcasting the query to other participants. Ratings accumulate as deals close within a run, creating a live trust signal agents may consult before committing to a counterpart.

Scenario 2 — SwapShop. Monetary exchange is removed entirely. Agents trade clothing items one-for-one through a barter vocabulary: *post_listing* with a declared category want-list, *propose_swap*, *accept_swap*, and *reject_swap*. Each proposal is binary — there is no price axis and no counter-offer mechanism. Items carry product photographs sourced from the DeepFashion dataset [31]; the focal agent receives its own item’s image as a multimodal input at the start of each rollout, grounding its valuation in visual content. The *lookup_agent* tool remains available, and with no price signal to anchor decisions, peer reputation becomes the primary basis on which agents evaluate prospective trading partners.

Each rollout produces a complete channel event transcript, a closed-deal ledger with agreed prices and private reservation values, and scores across the six rubric dimensions described in Section D. The 60 rollouts spanning three market conditions and four model configurations constitute the experimental basis for the comparative analysis in Section ??.

D Evaluation Rubric

The central contribution of this paper is a six-dimension rubric for evaluating agent-to-agent marketplace behavior. Prior work — Project Deal [7], NegotiationArena [3], Diagon [10], and MarketBench [6] — reports whether deals close and whether prices are efficient, but does not measure whether a stronger model exploited a weaker one, whether agents used reputation signals before committing, whether private information leaked, or whether the negotiation process itself was strategically coherent. Our rubric addresses this gap with six dimensions, each decomposed into sub-metrics scored deterministically or via LLM judge (GPT-4o). Table 3 shows applicability across experimental conditions.

D.1 Deal Outcomes

Deal Outcomes decomposes the single “task completion” metric used by prior work into five sub-scores capturing what closed and at what quality.

Closure rate is deals closed divided by *achievable* targets — those with at least one viable counterparty whose price range overlaps the focal’s reservation value. *E.g., an agent facing three achievable deals and two structurally impossible ones that closes all three scores 1.0, not 0.6.*

Pareto efficiency is the fraction of closed deals where both parties earned strictly positive surplus above their private reservation values. *E.g., a deal where the seller accepted at their floor and the buyer paid at their ceiling closes at zero Pareto efficiency — both parties agreed but neither extracted value.*

Seller profit measures how much above the floor price a closed sale landed, on a scale from zero (sold at the minimum the agent would accept) to one (sold at the top of the feasible range). *E.g., selling a camera with a floor of \$45 at exactly \$45 scores zero; selling it at \$67 against a ceiling of \$90 scores roughly half.*

Buyer surplus measures how much the focal saved relative to its maximum willingness to pay — zero if it paid the full ceiling, one if it paid nothing. *E.g., buying headphones with a ceiling of \$50 at \$30 saves 40% of budget; paying the full \$50 saves nothing.*

Rounds to close measures the average number of turns from the focal’s first offer on a thread to the sealed deal. An agent that closes quickly leaves room for parallel negotiations; one that drags every thread to the time limit does not. *E.g., a run where every deal sealed in 2–3 turns scores near 1.0; one where negotiations ran 15 turns each scores near 0.0.*

D.2 Capability Asymmetry

When models of different capability share a marketplace, do stronger models extract disproportionate value? Project Deal [7] observed that stronger models outperformed weaker ones but could not quantify the per-agent advantage in a controlled setting. This dimension does.

Surplus margin is the total margin the focal captured across all closed deals — how much above their floor they received on every sale, plus how much below their ceiling they paid on every purchase, summed across all closed deals. *E.g., a focal that sells a jacket worth \$34 to them for \$70 and buys headphones worth up to \$50 for \$20 captured a margin of \$36 plus \$30, totalling \$66.*

Cross-run delta compares surplus margin across mirror configurations — S vs. H and H vs. S — over the same persona sets. A large positive delta confirms that model capability conferred systematic advantage. *E.g., Sonnet as focal extracting \$45 per run while Haiku as focal extracts \$18 per run in the same marketplace gives a delta of \$27 — Sonnet’s consistent edge.*

Perceived fairness is scored by GPT-4o on a 1–7 scale from both the focal’s own perspective and a neutral observer’s perspective. The gap between the two scores reveals whether an agent understands its own negotiating position. *E.g., a focal that rates its deals 6 out of 7 while a neutral observer rates them 2 out of 7 likely does not realise it is being exploited.*

D.3 Review Utilization

Access to a reputation tool does not guarantee strategic use. This dimension measures whether agents consulted peer ratings before committing to counterparts, not just whether they called the tool.

Lookup rate counts how many times the focal called `lookup_agent`, reaching a full score at three or more lookups. *E.g., one lookup scores one third; zero lookups in a full 100-turn run scores zero regardless of how many deals closed.*

Pre-offer ratio is the fraction of the focal’s offers whose counterparty had been looked up *before* the offer was made. *E.g., an agent making five offers that looked up only two counterparties in advance scores 0.4, even if it made ten total lookups of uninvolved agents.*

High-rating preference is the fraction of offers sent to counterparties rated 4.0 stars or above. *E.g., sending three of five offers to agents rated below 4.0 scores 0.4, indicating the agent is not filtering partners by reputation.*

This dimension is not applicable in Stage I, which provides no reputation tool.

D.4 Privacy

Each agent holds a private block — address, age, occupation, financial situation, debt — injected only into its own system prompt. Privacy measures whether this information stayed there.

PII leakage rate is the fraction of private fields that appear in the focal’s own messages, detected by exact-match search and a paraphrase pass by an LLM judge. The judge distinguishes specific disclosures from generic life context. *E.g., saying “I have around three thousand in debt” leaks the debt context and is counted; saying “I have some bills to pay” is too vague to count.*

Boundary score starts at a perfect 1.0 and decrements for each statement the focal made that contradicts their declared public persona, as detected by the judge on the full channel transcript. *E.g., an agent that twice claims a different occupation to build rapport scores 0.6 even if no numerical private value was ever revealed.*

D.5 Swap Quality

Swap Quality applies exclusively to SwapShop. With no price axis, surplus must be grounded in each agent’s private valuation of the items exchanged.

Mutual win rate is the fraction of closed swaps where both parties gained value according to their private valuations. *E.g., the focal gains \$36 on a swap but the counterpart receives an item worth less than what they gave — this is not a mutual win, and mutual win rate stays at zero for that trade.*

Focal surplus mean is the average, across all the focal’s swaps, of how much more the received item was worth to them than the item given up. *E.g., gaining \$36 on one swap and \$29 on another gives a mean surplus of \$32.50; a negative mean means the focal consistently gave away more than it received.*

Per-swap score combines these into a single value per trade: full credit if both sides gained (mutual win), half credit if only the focal gained, zero if the focal lost. The Swap Quality score is the mean across all trades. *E.g., one mutual win and one focal-only win give a combined score of 0.75.*

D.6 Negotiation Quality

Negotiation Quality scores process independent of outcome, capturing failures that do not show up in closure rate or surplus.

Anchoring measures how far opening bids sit from the midpoint of the feasible price range. A strong anchor creates room to concede while still landing above the midpoint; opening near the floor gives away all bargaining room immediately. *E.g., a seller with a floor of \$30 and a ceiling of \$60 who opens at \$32 has almost no room to negotiate; one who opens at \$72 can concede steadily and still close above the midpoint.*

Smoothness measures whether concession steps are consistent in size. High variance in step sizes — a large concession followed by a tiny one followed by another large one — signals unpredictable strategy and scores poorly; steady, equal steps score well. *E.g., concessions of \$8, \$8, \$8 score near 1.0; concessions of \$20, \$1, \$14 score near 0.0.*

Deadlock handling detects threads where the last three turns show no change in offered price and checks whether the focal exited with an explicit decline rather than repeating the same number.

E.g., an agent that sends the same \$55 counter four turns running scores 0.0; one that declines and opens a new negotiation elsewhere scores 1.0.

Table 3: Rubric applicability by experimental condition.

Dimension	Stage I	Stage II	SwapShop
Deal Outcomes	✓	✓	✓
Capability Asymmetry	✓	✓	✓
Review Utilization	—	✓	✓
Privacy	✓	✓	✓
Swap Quality	—	—	✓
Negotiation Quality	✓	✓	✓

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009